

Hashing and Bloom filters key lookup comparison

names

March 25, 2020

Abstract

[illegible]

1 Introduction

This project is mainly aimed at learning different implementations of hashing based data structures and their effectiveness in applying a dictionary search problem.

A dictionary is used to maintain a set under insertion and deletion of elements while membership queries provide access to the data. The most efficient dictionaries, in theory, and practice, are based on hashing techniques.

To achieve different results and build a complete analysis, we have proposed searching with two types of dictionary data structures: hash tables and bloom filters.

The main performance parameters that we are going to study, analyze and contrast along this research work are: lookup time, build time, average probes and false positives. All the algorithms have been implemented in c++ and the same data has been applied to the execution of different versions of the same experiment.

2 Experiment pipeline and methodology

The experimentation process was a very tedious task since it required to take in account a lot of parameters. TODO(Add pipeline image)

2.1 Open Addressing

Open addressing is one of the two main methods of collision resolution in hash tables. The four main techniques are:

- **Linear probing:** its principle is accomplished using two values: a start index and a stepping value. When a collision occurs, the table is search sequentially for an empty slot adding repeatedly the second value to the first until either a free space is found or the entire table is traversed.
- **Quadratic probing:** its principle is accomplished using four values: a start index, a stepping value and two constant values. When a collision occurs, an arbitrary quadratic polynomial value -made by the combination of all the values but the start index- is added to the first one; the finish criteria is the same as the previous technique. The idea is to skip regions in the table with possible clusters.
- **Double hashing:** its principle is accomplished using two values: a start index, a stepping value and a third value. When a collision occurs, the combination of the stepping and the third value is added to the first one; the finish criteria is the same as the first technique. The idea is to tackle effectively clustering problems.
- **Cuckoo hashing:** its principle is accomplished using two values and either position is computed by the first or the second. When a collision occurs, the key in the position is replaced and the replaced key is assigned to the position given by the other value. If this new position is occupied, the procedure is repeated until the key is inserted or the iteration reaches an arbitrary value. The replacement behaviour may result in an infinite loop.

Technique	Function
LP	$h(k,i) = (h(k) + i) \bmod m$
QP	$h(k,i) = (h(k) + c_1i + c_2i^2) \bmod m$
DH	$h(k,i) = (h_1(k) + i * h_2(k)) \bmod m$
QH	$h(k) = h_1(k) \text{ or } h_2(k)$
Parameters	Meaning
k	key, value to insert or to search
i	stepping value
c_1, c_2	constant values
$h(), h_1(), h_2()$	hash functions

2.2 Separate Chaining

2.3 Bloom filters

Bloom filters [1] are a space efficient probabilistic data structure that can represent a dictionary. This data structures are space efficient because unlike hash tables they don't store the key itself but the a series of bits that determine if the value is in the dictionary or not. This bits are represented in an array or table. To insert a value into a bloom filter, it has to be hashed with k different hash functions, each hash

returns a specific position of the bits table that will be marked with a 1. Thus, if a value is in the dictionary it means that if we hash the key to be queried k times every position will be marked with a 1 bit. Otherwise if a bit is marked with a 0 it means that the value is not in the dictionary. This means that we if considered k as a constant we can query and insert always in constant time. On the other hand this causes some problems because it is possible that two keys k_1, k_2 have hash to the same k bits, and because we don't have a collision mechanism we would introduce a false positive rate when quering a specific value in the dictionary. This rate is directly dependent on the the number of hash functions that we use and the load factor of the table. To minimize [2] this false positive rate different methods have been found although you can never make it dissapear. Another problem that we encountered when developing bloom filters was that we needed a way to generate k different hash functions, although we found a very efficient and simple solution [4] based on creating a new hash function with two initial ones and an iterator.

$$g(Key)_i = h_1(Key) + i * h_2(Key) \quad (1)$$

3 Data Generation

The data used in each of the experiments consisted of a sequence of unsigned integers, the data was written in two different files. One file contained the keys to insert into the dictionary and the other the text to find in the dictionary.

Linear Congruential Method

The randomness of the keys was based on the Linear congruential method (LCM) [5]. We needed to have experiments as reliable as possible thus we needed a way to control the data generation, that is why we used this method. The LCM is one of the oldest and best-known random number generators, it is also very easy to implement. This method generates a cyclic sequence of numbers based on an initial seed. Given a seed X_0 it is possible to generate the next number in the sequence

$$X_i = (X_{i-1} * a + c) \bmod m \quad (2)$$

where a is called the multiplier, c the incrementer and m the modulus. The size of the sequence depends on the parameters used, independently on the initial seed. We used a combination of parameters that yield the maximum size sequence, that is parameters that yield m numbers. This ensured us that that there where never repeated numbers in either of the two files since the size of m was sufficiently big.

4 Experiment parameters

An experiment is purely defined by the subject, represented by the dictionary to test, and the parameters in each the dictionary was specified and tested.

- **Size of n :** The keys file contained n numbers and the text contained $2*n+p*n$ where p is the percentage of keys that are already inserted in the dictionary. We had a lot of discussion in which n used and we tried different combinations

but they yielded very similar results so we finally chosed an $n = 10e^7$. That happens to be 10 million keys that where inserted into the dictionary, although the table size was much bigger. We tried to used a bigger size of n but we encountered the size of the files we created where extraordinary big, we tried to divide the files but that wasn't of much use since we stilled needed to save all the numbers in the hash tables. When we tried to to this the linux oom-killer killed to process because of the out of memory that was caused. Nevertheless other researches [3] used similar sizes of n so this size wasn't really an issue.

- **Number of repetitions Q :** Each experiment was defined by a series of parameters to get the best results and avoid overfitting of the data to a single sample we repeated each experiment with $Q = 100$ that is we performed the same experiment 100 times and averaged the results.
- **Load factor α :** The load factor is one of the variables that really helps into comparing each type of dictionary implementation, we used in total 9 different load factors ranging from 0.10 to 0.9 with an offset of 0.10. We could have subdivided the interval with a lower offset but after experimenting we found that there wasn't really a substantial difference.
- **Key percentage p :** The key percentage represents the percentage of keys that are in the keys (inserted in the dictionary) that are in the search file (the second file where the lookups are done). This variable did not have a big impact in the results since we average the lokuup time of each key individually and even if we chosed a really low key percentage we had a big enough Q that it didn't matter. Thus we chosed $p = 0.5$;
- **Number of hash functions k :** This was a difficult parameter to choose since it had to be chosen optimally because if had a very big k then the number of false positives might diminish, but only if the size of the table was big enough, on the other hand if we chosed a small value of k the table would fill more slowly but it was more probable to find to keys with the same bits mask. To decide into what to choose we simply tested different values of k with different results.
- **Hash functions:** The hash functions we chosed where not the purpose of this experiment although we tried to choose.
- **Random seed:** The seed number is really a determining factor into how the dictionaries will perform since different seeds yield to different keys and conscuently differnet hash outputs. That is why for each experiment we chosed the same initial seed, although the next seed was chosen at random based on the initial seed, this enabled reproducability. In total a number of Q different seeds were selected, because the results of the Q repetitions where average this enabled us to remove any outliers that could happend choosing a specific seed value.

5 Results

All the results that were obtained are specifically compared to the load factor of the table since it is a determining parameter in the efficiency of each dictionary.

Speed performance

To compare the speed...

Probes metric

In this section we compare the average probes that were introduced in each search query...

False positives

This section is specifically centred in the false positive rate of the bloom filters...




Figure 1: False positive rate comparison

6 Discussion

WE LEARNED THAT....

References

- [1] Burton Howard Bloom. “Space/Time Tradeoffs in Hash Coding with Allowable Errors.” In: *Communications of the ACM* (1970).
- [2] Andrei Broder and Michael Mitzenmacher. “Network Applications of Bloom Filters: A Survey”. In: *Internet Mathematics* (2004).
- [3] Bernhard Grill. “A Survey on Efficient Hashing Techniques in Software Configuration Management”. In: (2014).
- [4] Adam Kirsch and Michael Mitzenmacher. “Less Hashing, Same Performance: Building a Better Bloom Filter”. In: *Wiley InterScience* (2007).
- [5] Donald Knuth. *The art of computer programming, Volume 2*. Addison Wesley, 2011.
- [6] Donald Knuth. *The art of computer programming, Volume 3*. Addison Wesley, 2011.
- [7] Professor Clark F. Olson and Carol Zander. “Hashing and Hash Tables.” In: *University of Washington Bothell’s web site*. (2017).
- [8] Rasmus Pagh and Flemming Friche Rodler. “Cuckoo Hashing.” In: *Brics: Basic Research in Computer Science*. (2001).
- [9] Abubakar Sulaiman Gezawa Abdurra’uf Garba Saifullahi Aminu Bello Ahmed Mukhtar Liman and Abubakar Ado. “Comparative Analysis of Linear Probing, Quadratic Probing and Double Hashing Techniques for Resolving Collusion in a Hash Table.” In: *International Journal of Scientific and Engineering Research*. (2014).